

Parallelism (PAR)

Mastering your task decomposition strategies:
going some steps further

Eduard Ayguadé, Gladys Utrera and Daniel Jiménez-González

Computer Architecture Department
Universitat Politècnica de Catalunya

Course 2025/26 (Spring semester)

Parallelism (PAR)

taskwait vs. taskgroup

```
#pragma omp task {}          // T1
#pragma omp task             // T2
{
    #pragma omp task {}     // T3
}
#pragma omp task {}         // T4

#pragma omp taskwait
// Only T1, T2 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}         // T5
```

```
#pragma omp task {}           // T1
#pragma omp taskgroup
{
    #pragma omp task           // T2
    {
        #pragma omp task {}    // T3
    }
    #pragma omp task {}        // T4
}
// Only T2, T3 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}           // T5
```

Outline

Video lesson 4

Task Decomposition

Iterative vs Recursive

Task generation control

Iterative task decompositions

Recursive task decompositions

Exploratory recursive problems

Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

Iterative vs Recursive

Task generation control

Iterative task decompositions

Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems

Task Decompostion

- ▶ **Iterative:** A task may be a "code block", a procedure invocation, a set of loop iterations or full loops.
- ▶ **Recursive:** Tasks found in divide-and-conquer problems and other recursive problems. In particular, we distinguish ...
 - ▶ Leaf: each base case of the sequential recursive code is defined as a task
 - ▶ Tree: each recursive call is defined as a task

Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

Iterative vs Recursive

Task generation control

Iterative task decompositions

Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems

- ▶ In iterative task decomposition strategies one can control task granularity by setting the number of iterations executed by each task
- ▶ In recursive task decomposition strategies one can control task granularity by controlling recursion levels where tasks are generated (**cut-off control**)
 - ▶ after certain number of recursive calls (static control)
 - ▶ when the size of the vector is too small (static control)
 - ▶ when there are sufficient tasks pending to be executed (dynamic control)
 - ▶ ...

Outline

Video lesson 4

Task Decomposition

Iterative vs Recursive

Task generation control


Iterative task decompositions

Recursive task decompositions

Exploratory recursive problems

Schedule clause in for work-sharing (optional)

Different options to assign chunks of iterations to each implicit task through the `schedule` clause



Outline

Video lesson 4

Task creation and synchronization (Labs summary)

Task Decomposition

Iterative vs Recursive

Task generation control

Iterative task decompositions


Recursive task decompositions

Reducing overheads and serialization due to synchronization

Exploratory recursive problems


Recursive task decomposition: divide-and-conquer (2)

N=1024, MIN.SIZE=64



Recursive task decomposition: leaf strategy (1)

A task corresponds with each invocation of `dot_product` once the recursive invocations stop



- Sequential generation of tasks

Leaf strategy: where is the task synchronization? (1)

```

#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}

```

- ▶ Where is the task synchronization?
- ▶ Are there nested tasks?

Leaf strategy: where is the task synchronization? (2)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        // Now we need the result here.
        ....
    }
}
```

- What kind of synchronization should we use? Where?

Leaf strategy: where is the task synchronization? (3)

```

#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);


void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        #pragma omp taskwait
        // Now we need the result here.
        ....
    }
}

```


How to control task granularity in leaf strategy (1)

Leaf parallelization with **depth recursion control**




How to control task granularity in leaf strategy (2)

Leaf strategy with **depth recursion control**

```
#define CUTOFF 2
...
void rec_dot_product(int *A, int *B, int n, int depth) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth == CUTOFF)
            #pragma omp task
            {
                rec_dot_product(A, B, n2, depth+1);
                rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            }
        else {
            rec_dot_product(A, B, n2, depth+1);
            rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else // if recursion finished, need to check if task has been generated
        if (depth <= CUTOFF)
            #pragma omp task
            dot_product(A, B, n);
        else
            dot_product(A, B, n);
}
...
```

Recursive task decomposition: tree strategy (1)

A task corresponds with each invocation of `rec_dot_product`



- ▶ Parallel generation of tasks
- ▶ Granularity: some tasks simply generate new tasks

Recursive task decomposition: different sequential code ...

```

int dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    return(tmp);
}

int rec_dot_product(int *A, int *B, int n) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        tmp1 = rec_dot_product(A, B, n2);
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);
    } else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}

void main() {
    result = rec_dot_product(a, b, N);
}

```

Recursive task decomposition: tree strategy (2)

```

int dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    return(tmp);
}


int rec_dot_product(int *A, int *B, int n) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task shared(tmp1) // firstprivate(A, B, n, n2) by default
        tmp1 = rec_dot_product(A, B, n2);
        #pragma omp task shared(tmp2) // firstprivate(A, B, n, n2) by default
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);
        #pragma omp taskwait
    } else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    result = rec_dot_product(a, b, N);
}

```

How to control task granularity in tree strategy (1)

Tree strategy with **depth recursion control**



How to control task granularity in tree strategy (2)

Tree strategy with **depth recursion control**

```
#define N 1024
#define MIN_SIZE 64
#define CUTOFF 3

int rec_dot_product(int *A, int *B, int n, int depth) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth < CUTOFF) {
            #pragma omp task shared(tmp1)
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            #pragma omp task shared(tmp2)
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            #pragma omp taskwait
        } else {
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else tmp = dot_product(A, B, n);
    return(tmp1+tmp2);
}
```

OpenMP support for cut-off

- ▶ `final` clause: If the expression of a `final` clause evaluates to *true* the generated task and **all of its descendent tasks** will be final. The execution of a final task is sequentially **included** in the generating task (but the task is still generated)
- ▶ `omp_in_final()` intrinsic function: it returns true when executed in a final task region; otherwise, it returns false.

Tree strategy: where is the task synchronization? (3)

```
int result = 0;
void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task
        rec_dot_product(A, B, n2);
        #pragma omp task
        rec_dot_product(A+n2, B+n2, n-n2);
    } else dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp taskgroup
        {
            rec_dot_product(a, b, N);
        }
        // Now we need the result here..
        ....
    }
}
```


Avoiding task barriers: task dependences (1)

- ▶ The OpenMP runtime detects dependences between sibling tasks (i.e. from the same parent task) through the specification of the directionality for the variables used in the tasks

```
#pragma omp task [depend (in : var_list)]
                  [depend (out : var_list)]
                  [depend (inout : var_list)]
```

Task dependences are derived from the directionality type (`in`, `out` or `inout`) and its items in `var_list`; this list may include array sections (e.g. `v[0:n]`)

Avoiding task barriers: task dependences (2)

- ▶ **in specifier:** the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `out` or `inout` list ¹
- ▶ **out and inout specifier:** the generated task will be a dependent task of all previously generated sibling tasks that reference at least one of the list items in an `in`, `out`, or `inout` list


Types of dependences:

- ▶ *read-after-write*: caused by matched out – in
- ▶ *write-after-read*: caused by matched in – out
- ▶ *write-after-write*: caused by matched out – out

¹ Note: if a list item is an array section, the matching should occur with an identically defined array section.

Serialisation caused by task barriers (1)

Given a TDG to implement with the OpenMP tasking model:




```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp task
foo3()
#pragma omp taskwait
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```


```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp taskwait
#pragma omp task
foo3()
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```




Serialisation caused by task barriers (2)

```
#pragma omp task
{
  #pragma omp task
  foo1()
  #pragma omp task
  foo2()
  #pragma omp taskwait
  #pragma omp task
  foo4()
  #pragma omp taskwait
}
#pragma omp task
foo3()
#pragma omp taskwait
#pragma omp task
foo5()
```



```
#pragma omp task
foo3()
#pragma omp taskgroup
{
  #pragma omp task
  foo1()
  #pragma omp task
  foo2()
}
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
```



```
#pragma omp task depend(out: a)
foo1()
#pragma omp task depend(out: b)
foo2()
#pragma omp task depend(out: c)
foo3()
#pragma omp task depend(in: a, b)
                        depend(out: d)
foo4()
#pragma omp task depend(in: c, d)
foo5()
```



- ▶ `taskwait` with `depend` clause: instead of waiting for all child tasks to complete execution, it only waits for the predecessor child tasks according to the `in`, `out` and `inout` specifiers

```
#pragma omp taskwait
printf("final values for x=%d ; y=%d\n", x, y);
```

Additional functionalities (2) (optional)

- ▶ An iterator can be used in the depend clause, expanding to multiple values in the specifier they appear

```
for (i = 0; i < n; ++i)
    if (i%2) {
        #pragma omp task depend(out: v[i])
        compute_element(&v[i], i);
    }


#pragma omp task depend(iterator(it = 0:n), in: v[it])
                        // could also be depend(iterator(it = 1:n:2), in: v[it])
odd = sum_odd_elements(v, n);

even = sum_even_elements(v, n);
```

Note: this is not equivalent to the use of an array section in the `in` specifier (i.e. `depend(in:v[0:n])`), why not?

Additional functionalities (3) (optional)

- `mutexinoutset` specifier: equivalent to `inout` but all dependent tasks can be executed in any order, one after the other



Red dependence expressed with `depend(inout:x)` (top temporal diagram) or with `depend(mutexinoutset:x)` (bottom temporal diagram). Observe that tasks can be executed in any order, but only one at a time.

Cancellation points in OpenMP (optional)

Tasks induced by exploratory decomposition can be terminated before finishing as soon as the desired solution is found

- ▶ `#pragma omp cancel [parallel | taskgroup]`: this directive activates the cancellation of the enclosing `[parallel | taskgroup]` region. The thread that finds the directive finishes its execution; the other threads continue their execution as normal.
- ▶ `#pragma omp cancellation point [parallel | taskgroup]`: introduces a point to check if cancellation has been activated. When found by a thread, if the enclosing `[parallel | taskgroup]` region has been already cancelled, then it finishes its execution.

Cancellation points in OpenMP: very simple example (optional)

```
#pragma omp taskgroup
for (i=0; i<1000; i=i+100)
    #pragma omp task firstprivate(i) private(j)
    {
        for (j=i; j<i+100; j++) {
            if (do_computation(j) == 0) {
                #pragma omp cancel taskgroup
            }
            #pragma omp cancellation point taskgroup
        }
    }
```

The first task with 0 as a result of `do_computation` will finalise the execution of all the tasks in the `taskgroup`

Protecting task interactions in OpenMP (Labs summary)

Two mechanisms:

1. Atomic accesses: mechanism to guarantee atomicity in load/store instructions

```
#pragma omp atomic [update | read | write]  
    expression
```

- ▶ Atomic updates: `x += 1`, `x = x - foo()`, `x[index[i]]++`
- ▶ Atomic reads: `value = *p`
- ▶ Atomic writes: `*p = value`

Protecting task interactions in OpenMP (Labs summary)

Two mechanisms:

- 2. Mutual exclusion: mechanism to ensure that only one task at a time executes the code within a critical section
 - ▶ `critical` pragma: a thread waits at the beginning of a critical region until no other thread is executing a critical region (anywhere in the program)
 - ▶ `critical(name)` pragma: the `name` allows the programmer to differentiate disjoint sets of critical sections (`name` is a label, not a program variable)
 - ▶ `omp_lock_t` OpenMP intrinsics and low-level synchronization primitives (next in this chapter)

Reducing task interactions: overhead (1)

Reductions: replicate key data structures and locally working with these local structures; when appropriate, locally replicated data structures are combined into the final global result

```
int result = 0;
// Assume this function is instantiated as a task
void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        #pragma omp atomic
        result += A[i] * B[i];
}
```

could be easily transformed into

```
void dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++)
        tmp += A[i] * B[i];

    #pragma omp atomic
    result += tmp;
}
```

Reducing task interactions: overhead (2)

Specifying reduction operations in explicit tasks generated with either `task`:

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp taskgroup task_reduction(+: sum)
    for (i=0; i< SIZE; i++)
        #pragma omp task firstprivate(i) in_reduction(+: sum)
        sum += X[i];
}
```

or `taskloop` (possible because this loop is a countable):

```
#pragma omp parallel
#pragma omp single
{
    // implicit taskgroup in taskloop construct
    #pragma omp taskloop reduction(+: sum)
    for (i=0; i< SIZE; i++)
        sum += X[i];
}
```

Low-level synchronization functions using *locks*

Locks: special variables that live in memory with two basic operations:

- ▶ **Acquire:** while a thread has the lock, nobody else gets it; this allows the thread to do its work in private, not bothered by other threads
- ▶ **Release:** allow other threads to acquire the lock and do their work (one at a time) in private

Type definition and intrinsics:


```
void omp_init_lock(omp_lock_t *lock)
void omp_destroy_lock(omp_lock_t *lock)

void omp_set_lock(omp_lock_t *lock)
void omp_unset_lock(omp_lock_t *lock)

int omp_test_lock(omp_lock_t *lock)
```

Reducing task interactions: serialization (1)

Example: inserting elements in hash table defined as a collection of linked lists



```
typedef struct {  
    int data;  
    element *next;  
} element;  
  
int dataTable[SIZE_TABLE];  
element * HashTable[SIZE_HASH];  
  
for (i = 0; i < SIZE_TABLE; i++) {  
    int index = hash_function (dataTable[i], SIZE_HASH);  
    insert_element (dataTable[i], index, HashTable);  
}
```

Reducing task interactions: serialization (2)

Easily parallelizable using an iterative task decomposition using `taskloop`. However ...

- ▶ ... updates to the list in any particular slot must be protected to prevent a race condition

```
typedef struct {
    int data;
    element *next;
} element;

int dataTable[SIZE_TABLE];
element * HashTable[SIZE_HASH];

#pragma omp taskloop
for (i = 0; i < elements; i++) {
    int index = hash_function (dataTable[i], SIZE_HASH);
    #pragma omp critical // atomic not possible here
    insert_element (dataTable[i], index, HashTable);
}
```

- ▶ Serialization in the insertion of elements

Reducing task interactions: serialization (3)

Associate a lock variable with each slot in the hash table, protecting the chain of elements in an slot

```
omp_lock_t hash_lock[SIZE_HASH];

#pragma omp parallel
#pragma omp single
{
    for (i = 0; i < SIZE_HASH; i++) omp_init_lock(&hash_lock[i]);

    #pragma omp taskloop
    for (i = 0; i < SIZE_TABLE; i++) {
        int index = hash_function (dataTable[i], SIZE_HASH);
        omp_set_lock (&hash_lock[index]);
        insert_element (dataTable[i], index, HashTable);
        omp_unset_lock (&hash_lock[index]);
    }

    for (i = 0; i < SIZE_HASH; i++) omp_destroy_lock(&hash_lock[i]);
}
```

Threads may be inserting elements into the hash table in parallel, as long as these elements hash to different slots

How would you address the N-queens problem? (1)

```


char *a; // Solution being explored
int sol_count = 0; // Total number of solutions found
int size = 8; // board size

void nqueens(int n, int j, char *a) {
    if (j == n) sol_count += 1;
    else
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a)) nqueens(n, j + 1, a);
        }
}

int main() {
    a = alloca(size * sizeof(char));
    nqueens(size, 0, a);
}


```

a = [0, 6, 3, 5, 7, 1, 4, 2]



How would you address the N-queens problem? (2)

For a 4x4 board, the recursion tree would be ...



How would you address the N-queens problem? (3)

```

void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a))
                #pragma omp task                                // all firstprivate by default
                nqueens(n, j + 1, a);
        }
    // Do we need to insert a task barrier at this point?
}

int main() {
    a = alloca(size * sizeof(char));
    #pragma omp parallel
    #pragma omp single
    nqueens(size, 0, a);
}

```

Do we need a new board for each task to be able to explore its own path? Is the implicit `firstprivate(a)` enough?

How would you address the N-queens problem? (4)

A new board has to be allocated if the path is explored as a task

```
void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else {
        // try each possible position for queen <j>
        for ( int i=0 ; i < n ; i++ ) {
            a[j] = (char) i;
            if (ok(j + 1, a)) {
                // allocate a temporary array and copy <a> into it
                char * b = alloca(n * sizeof(char));
                memcpy(b, a, (j + 1) * sizeof(char));
                #pragma omp task                                // all firstprivate by default
                nqueens(n, j + 1, b);
            }
            #pragma omp taskwait
        }
    }
}
```

Important: firstprivate(b) (implicit for new board) captures the pointer to b, not the whole vector b

Adding cut-off to N-queens (recursion level)


```

void nqueens(int n, int j, char *a) {
    if (j == n)
        #pragma omp atomic
        sol_count += 1;
    else
        // try each possible position for queen <j>
        if (!omp_in_final()) {
            for ( int i=0 ; i < n ; i++ ) {
                a[j] = (char) i;
                if (ok(j + 1, a))
                    // allocate a temporary array and copy <a> into it
                    char * b = alloca(n * sizeof(char));
                    memcpy(b, a, (j + 1) * sizeof(char));
                    #pragma omp task final(j>CUT_OFF)
                    nqueens(n, j + 1, b);
            }
            #pragma omp taskwait
        } else
            for ( int i=0 ; i < n ; i++ ) {
                a[j] = (char) i;
                if (ok(j + 1, a)) nqueens(n, j + 1, a);
            }
}

```


Another example: 15-puzzle (without code) ... (optional)

The solution to a 15-puzzle (a tile puzzle). Possible movements of the empty cell: UP, RIGHT, LEFT and DOWN. Here we show a series of moves that transform a given initial state to the desired final state:



Another example: 15-puzzle (without code) ... (optional)


The state space can be explored by generating various successor states of the current state and to view them as independent tasks



R: repeated configuration

Another example: 15-puzzle (without code) ... (optional)

Anomalous speed-ups of the parallel formulation of the problem:
the speed-up depends on where the solution is found ...



- ▶ Left: $T_1 = 2 \times m + 1$ and $T_4 = 1$, therefore ... $S_4 = 2 \times m + 1$
- ▶ Right: $T_1 = m$ and $T_4 = m$, therefore ... $S_4 = 1$

And the parallel efficiency (i.e. how well used are processors)?
Observe that on the right three processors waste their computation

